

Quick Guide to Managing Logs Efficiently

Optimize Log Retention: How IRQ and RwIQ Fit Your Strategy



Govern



Optimize



Escalate



Iterate



Govern Access

Action: Apply IAM restrictions to limit query access.

Role Guardrails



Contractors

RwIQ only, scoped buckets/time



Help Desk L1

Pre-made Notebooks & Dashboards
Simple analysis through logs app



Power Users

IRQ with alerts



SRE/Platform

Full access with audit

Guardrails Checklist

☐	Ingest - Are all log sources defined?
☐	Access Control - Who gets access to what data?
☐	Query - What is the DQL proficiency of users querying the data?
☐	Retention - How long do we have to keep this data?
☐	OpenPipeline® - Do we need to manipulate or process the data we are bringing in further?
☐	Future Actions - What is the plan for using this data in Dynatrace?

IAM Condition (sample): "storage:query-consumption": "INCLUDED"

Query Scope Guardrail (Example)

- If it is set to **INCLUDED**, executed queries return data only from the configured included queries timeframe of buckets from which RwIQ is set.
- If set to **ON_DEMAND**, it returns data from the entire timeframe.
- For **any other value**, the IAM statement is ignored. The included timeframe must be configured per bucket via the Storage Management API.



Spot Inefficient Queries

Action: Identify inefficient DQL and reduce scanned GiB.

Watch For

- ⚠ Wide time ranges → narrow to 24h
- ⚠ No bucket filter → target likely buckets
- ⚠ No filters → use filters + timeframe + scan limit
- ⚠ High-cardinality group-bys → summarize

Best Practice

- 🏆 Scope narrowly first (start with 24h + 1-2 buckets)
- 🏆 Filter, set timeframes, set scanLimits
- 🏆 Aggregate → drill (low-cardinality first)
- 🏆 Use guardrails

Narrow the Query Time Range

Example:

For better performance, limit the analysis time range

Code: fetch bizevents, from:-10m



Respond and Escalate

Action: Handle RwIQ overages or breaches and throttling calmly.

Immediate Steps

- Pause non-critical queries/jobs
- Narrow time windows; enforce RwIQ for high-volume buckets
- Communicate ETA clearly
- Apply optimizations; revisit included-query days via Estimator

Escalation Package

- Bucket config + retention settings
- Query logs (top offenders)
- ScannedGiB trend (include spike window)
- Estimator expected vs actual
- IAM mappings (IRQ/RwIQ roles)
- Actions taken

Policy Reminder: 15x Rule



• **Internal CTA:** #help-log-pb-grail / Licensing

• **Customer CTA:** Support/IAT route



Monitor and Iterate

Action: Track outcomes and adjust

Metrics to Watch



Ingestion volume



Daily scannedGiB



Query patterns



Estimator alignment



IAM exceptions

Measure

Adjust

Share